

BM 402 Bilgisayar Ağları (Computer Networks)

M.Ali Akcayol
Gazi Üniversitesi
Bilgisayar Mühendisliği Bölümü

Not: Bu dersin sunumları, ders kitabının yazarları James F. Kurose ve Keith W. Ross tarafından sağlanan sunumlar üzerinde değişiklik yapılarak hazırlanmıştır.

Transport Layer

Amaçlar :

- Transport layer hizmetlerinin anlaşılması:
 - multiplexing/demultiplexing
 - reliable data transfer
 - flow control
 - congestion control
- Internette transport layer protokolleri:
 - UDP: connectionless transport
 - TCP: connection-oriented transport
 - TCP congestion control


Ders konuları

- **Transport-layer hizmetleri**
- Multiplexing ve demultiplexing
- Connectionless transport: UDP
- Reliable data transferin prensipleri
- Connection-oriented transport: TCP
 - segment yapısı
 - reliable data transfer
 - flow control
 - connection management
- Tıkanıklık denetimi prensipleri
- TCP tıkanıklık denetimi

3/101

Transport hizmetleri ve protokoller

- Farklı hostlardaki prosesler arasında **mantıksal bağlantı (logical communication)** oluşturur
- transport protokolleri uç sistemlerde çalışır
 - Gönderen taraf: uygulama mesajlarını **segment**'lere böler ve network layer'a gönderir
 - Alıcı taraf: segmentleri birleştirir ve application layer'a gönderir
- Uygulamalar için birden fazla transport protokol bulunmaktadır
 - Internet: TCP ve UDP



4/101

Transport ve network layer

- **network layer:** hostlar arasında mantıksal bağlantı
- **transport layer:** prosesler arasında mantıksal bağlantı
 - Network layer hizmetlerine güvenir

Benzetim:


12 çocuk farklı yerde aynı evdeki 12 çocuğa mektup gönderiyor

- prosesler = çocuklar
- Uygulama mesajları = zarflardaki mektuplar
- hostlar = evler
- transport protokol = her evde mektup toplama ve dağıtma yapan kişiler
- network-layer protokol = posta hizmeti

5/101

Internet transport layer protokolleri

- Güvenilir, sıralı gönderim: TCP
 - congestion control
 - flow control
 - connection setup
- Güvenilir olmayan, sırasız gönderim: UDP
 - IP protokolünün best-effort özelliği kullanılır
- Sunulamayan hizmetler:
 - delay garanti
 - bandwidth garanti



6/101

Ders konuları

- Transport-layer hizmetleri
- **Multiplexing ve demultiplexing**
- Connectionless transport: UDP
- Reliable data transferin prensipleri
- Connection-oriented transport: TCP
 - segment yapısı
 - reliable data transfer
 - flow control
 - connection management
- Tıkanıklık denetimi prensipleri
- TCP tıkanıklık denetimi

7/101


Multiplexing/demultiplexing


Demultiplexing alıcı host'ta


Doğru soketten gelen segmentler alınır

Multiplexing gönderen host'ta

Birden çok soketten gelen veri toplanır, başlık eklenir (demultiplex için kullanılır)

 = soket


 = proses



8/101

Demultiplexing işlemi

- **host IP datagramlarını alır**
 - her datagram kaynak IP adresine ve hedef IP adresine sahiptir
 - her datagram 1 transport-layer segment taşır
 - eğer segment kaynak ve hedef port numarasına sahiptir (bazı uygulamalar için bilinen port numaraları atanır)
- **host IP adreslerini ve port numaralarını kullanarak segmenti uygun sokete yönlendirir.**



TCP/UDP segment format

9/101

Connectionless demultiplexing

- Port numaralarıyla soketler oluşturulur:

```
DatagramSocket mySocket1 = new  
    DatagramSocket(99111);  
DatagramSocket mySocket2 = new  
    DatagramSocket(99222);
```

- UDP soket bir ikiliyle tanımlanır:

(dest IP address, dest port number)

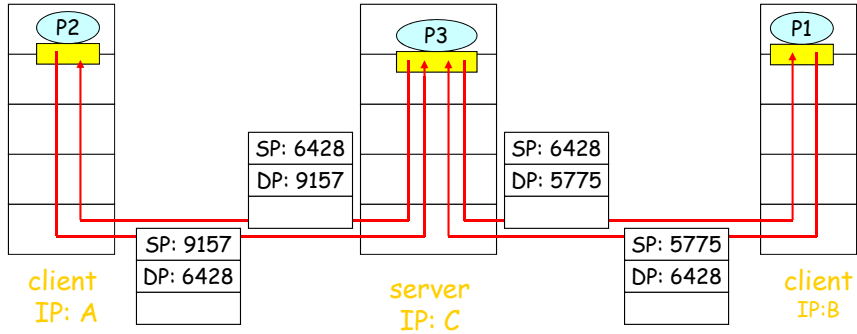
- Host UDP segment aldığı anda:

- Segmentteki hedef port numarası kontrol
- UDP segment port numarasıyla birlikte sokete yönlendirilir
- IP datagramlar aynı sokete kaynak ve hedef IP ve port numaralarıyla yönlendirilir

10/101

Connectionless demultiplexing - devam

```
DatagramSocket serverSocket = new DatagramSocket(6428);
```



SP dönüş adresini (return address) sağlar

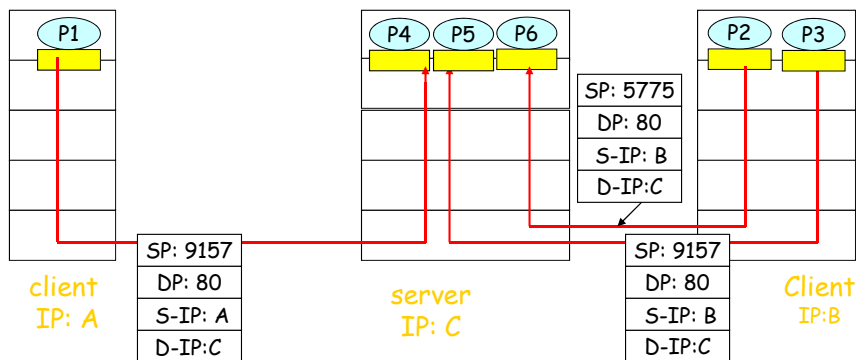
11/101

Connection-oriented demultiplexing

- TCP socket bir dörtlüyle tanımlanır:
 - source IP address
 - source port number
 - dest IP address
 - dest port number
- Alıcı host bu 4 değeri segmenti uygun sokete yönlendirmek için kullanır
- Server host eşzamanlı çok sayıda TCP socket destekler:
 - her socket kendi dörtlüsüyle tanımlanır
- Web sunucular her bağlanan client için farklı bir sokete sahiptir
 - non-persistent HTTP her istek için farklı bir sokete sahiptir

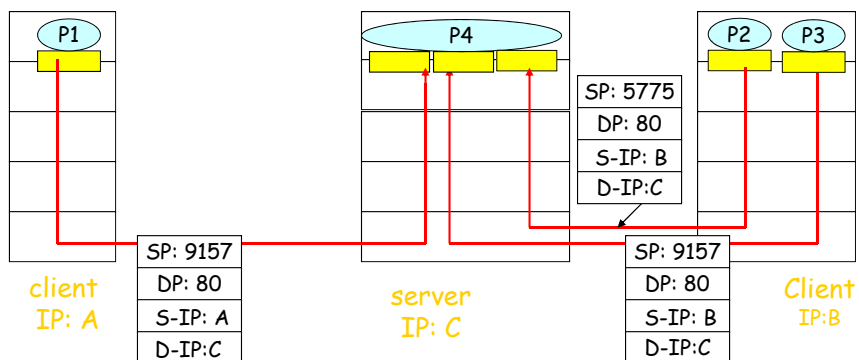
12/101

Connection-oriented demultiplexing - devam



13/101

Connection-oriented demux: Threaded Web Server



14/101

Ders konuları

- Transport-layer hizmetleri
- Multiplexing ve demultiplexing
- **Connectionless transport: UDP**
- Reliable data transferin prensipleri
- Connection-oriented transport: TCP
 - segment yapısı
 - reliable data transfer
 - flow control
 - connection management
- Tıkanıklık denetimi prensipleri
- TCP tıkanıklık denetimi

15/101

UDP: User Datagram Protocol [RFC 768]

- Güvenilir olmayan Internet transport protokolü
- “best effort” servis, UDP segmentleri:
 - Lost (kayıp olabilir)
 - Delivered out of order to app (uygulamaya sırasız gidebilir)
- **connectionless:**
 - UDP alıcı ve gönderici arasında handshaking yapılmaz
 - Her UDP segment diğerlerinden ayrı değerlendirilir

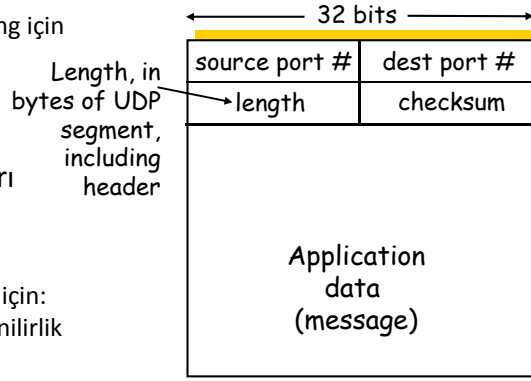
UDP niçin vardır?

- Bağlantı oluşturulmaz (gecikme olur)
- basittir: alıcı ve gönderici arasında bağlantı durumu yoktur
- Küçük segment başlığı
- Tıkanıklık denetimi yoktur: UDP istediği hızda veri gönderebilir.

16/101

UDP - devam

- Sıklıkla multimedia uygulamalarında streaming için kullanılır
 - loss tolerant
 - rate sensitive
- diğer UDP kullanımları
 - DNS
 - SNMP
- UDP ile güvenilir transfer için: application layer'da güvenilirlik eklenmelidir.



UDP segment format

17/101

UDP checksum

Amaç: iletilen segmentte hata algılama

Gönderici:

- segment içeriklerine 16-bit integer dizisi olarak bakılır
- checksum: segment içeriğinin 1 tümleyeni toplamı alınır
- Gönderici checksum değeri UDP checksum alanına yerleştirir

Alıcı:

- Alınan segmentte checksum hesaplanır
- hesaplanan checksum gelen checksum değeriyle karşılaştırılır:
 - HAYIR – hata var
 - EVET – hata yok. *Yinede hiçbir şekilde hata yok denilebilirmi?*

18/101

Internet checksum örnek

- Sayıları toplarken en soldaki bitlerdeki taşma sonuca eklenir
- Örnek: iki 16-bit integer toplanırsa

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

19/101

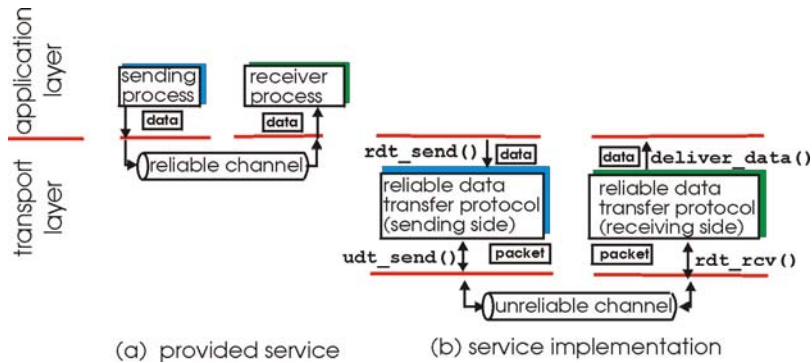
Ders konuları

- Transport-layer hizmetleri
- Multiplexing ve demultiplexing
- Connectionless transport: UDP
- **Reliable data transferin prensipleri**
- Connection-oriented transport: TCP
 - segment yapısı
 - reliable data transfer
 - flow control
 - connection management
- Tıkanıklık denetimi prensipleri
- TCP tıkanıklık denetimi

20/101

Reliable data transfer prensipleri

- Application, transport ve data link layer'da önemlidir
- Ağlarda en öncelikli 10 konunun içindedir.



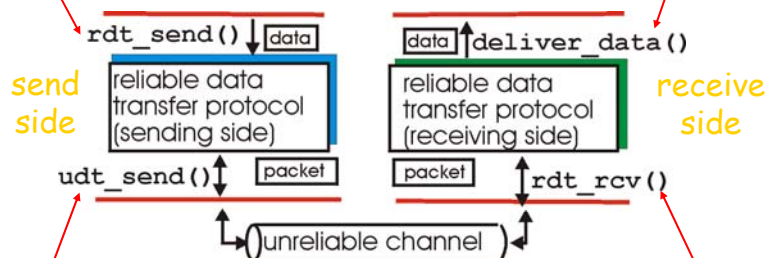
- Reliable data transfer (rdt) protokol daha karmaşıktır.

21/101

Reliable data transfer

rdt_send(): called from above, (e.g., by app.). Passed data to deliver to receiver upper layer

deliver_data(): called by rdt to deliver data to upper



udt_send(): called by rdt, to transfer packet over unreliable channel to receiver


rdt_rcv(): called when packet arrives on rcv-side of channel

22/101

Reliable data transfer

- reliable data transfer (rdt) protokolün gönderici ve alıcı taraflarını geliştirelim
- Data transferin tek yönlü (unidirectional) olduğunu düşünürsek
 - Ancak kontrol bilgisi iki yönlü gitmektedir
- Sonlu durum makineleriyle (finite state machines-FSM) modellenebilir


state: bir sonraki durum bu durumdayken oluşan sonraki olayla belirlenir



23/101

Rdt 1.0: reliable kanal kullanarak reliable transfer

- Altyapıdaki kanal tümüyle güvenilirdir
 - Bit hatası yoktur
 - Kayıp paket yoktur
- Gönderici ve alıcı için FSM:
 - Gönderici kanala veriyi gönderir
 - Alıcı kanaldan gelen veriyi okur



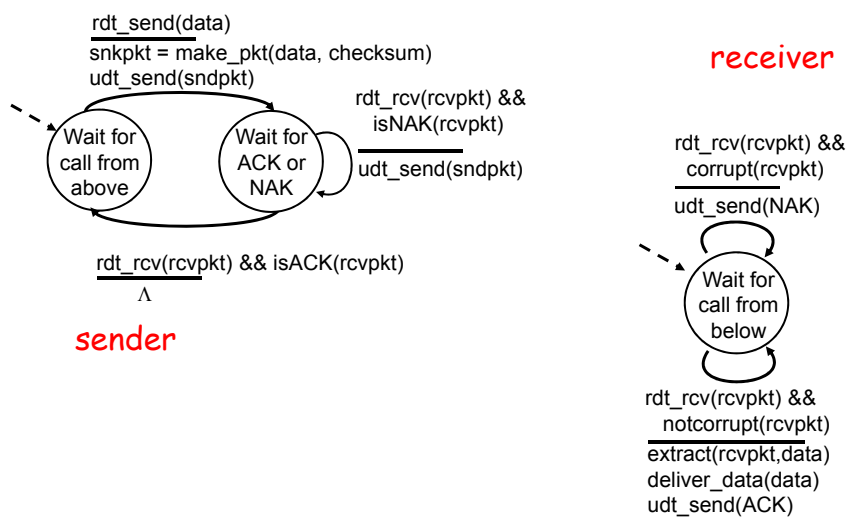
24/101

rdt 2.0: bit hatası olan kanal ile çalışma

- Kanalda paket içindeki bitlerde bozulma olabilir
 - Bit hatalarını kontrol etmek için checksum kullanılır
- Hatalar nasıl düzeltilir ?
 - **acknowledgements (ACKs)**: alıcı göndericiye aldığı paketin hatasız olduğunu iletir.
 - **negative acknowledgements (NAKs)**: alıcı göndericiye aldığı paketin hatalı olduğunu bildirir
 - Gönderici NAK ile bildirilen paketi tekrar gönderir
- **rdt2.0** daki yenilikler(**rdt1.0 a göre**):
 - Hata denetimi
 - Alıcı geri bildirimi: kontrol mesajları (ACK, NAK) alıcı->gönderici

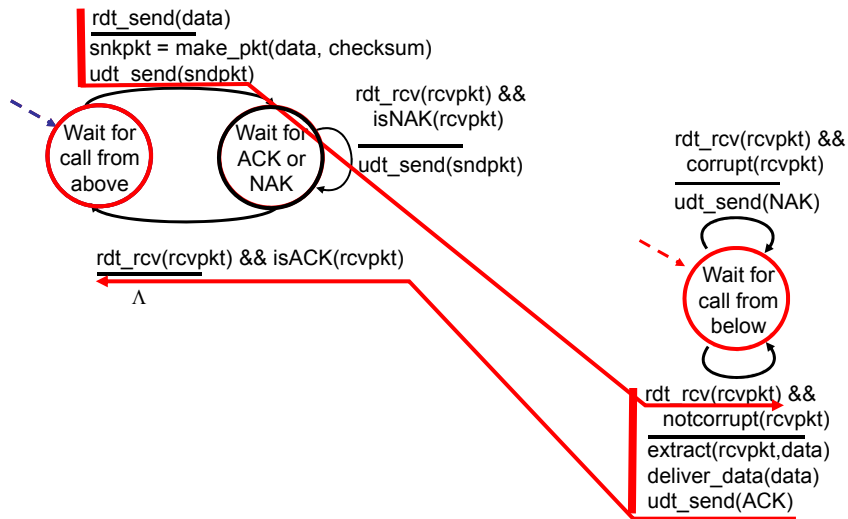
25/101

rdt2.0: FSM özellikleri



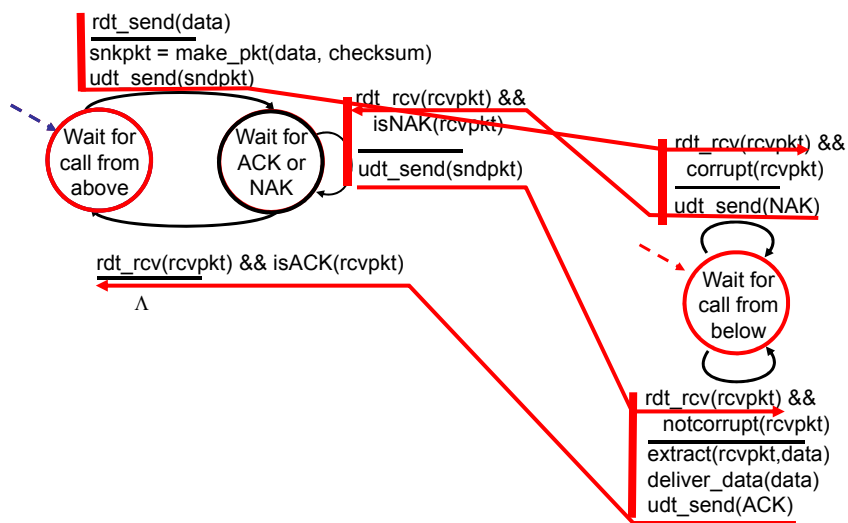
26/101

rdt2.0: hata olmadığı zaman çalışma



27/101

rdt2.0: hata durumu



28/101

rdt2.0 da karşılaşılan problemler

ACK/NAK bozulursa ?

- Gönderici alıcıda ne olduğunu bilemez
- Retransmit yapılmaz: duplicate olabilir

Duplicate'lerin seçilmesi:

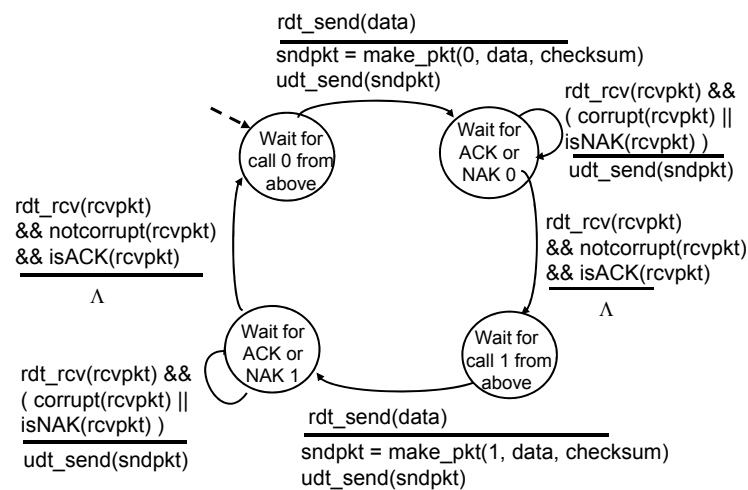
- Gönderici her pakete *sequence number* ekler
- Gönderici mevcut paketi retransmit yapar ACK/NAK bozulursa
- alıcı duplicate paketleri atar

stop and wait

Gönderici bir paket gönderir,
Alıcıdan cevap bekler

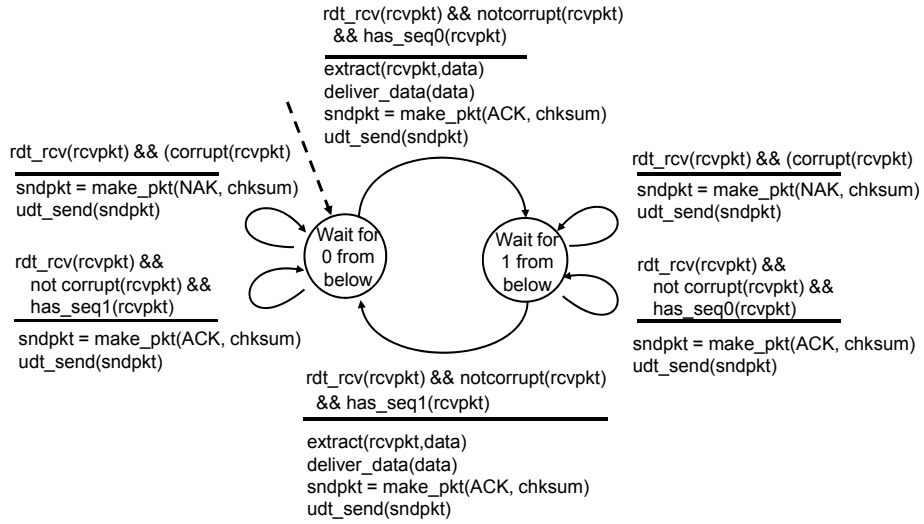
29/101

rdt2.1: gönderici bozulan ACK/NAK ları belirler



30/101

rdt2.1: gönderici bozulan ACK/NAK ları belirler



31/101

rdt2.1: değerlendirme

Gönderici:

- seq # pakete eklenir
- İki seq. no (0,1) yeterlidir.
- alınan ACK/NAK paketin bozuk olup olmadığı kontrol edilir

Alıcı:

- Gelen paket çiftmi kontrol edilir
 - Bulunulan durum gelen paket için seq. no, 0 veya 1 olarak bekler

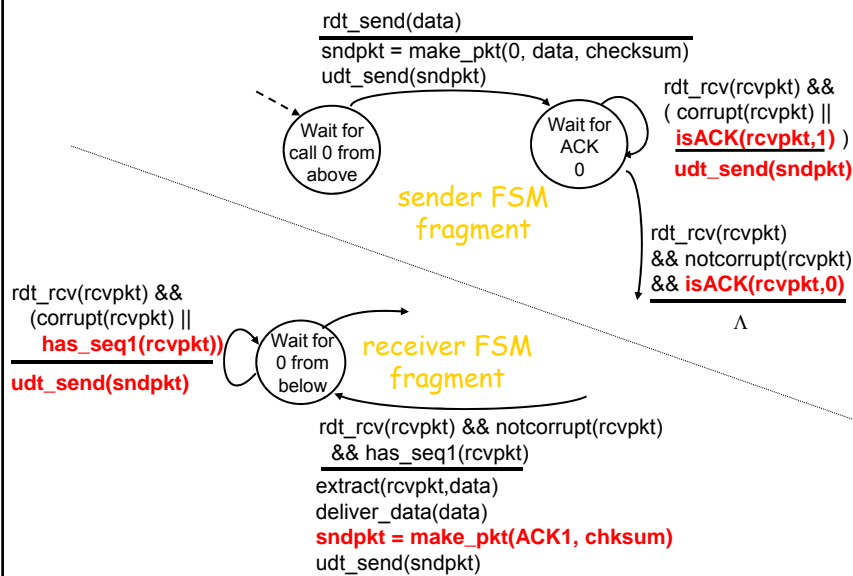
32/101

rdt2.2: NAK kullanılmayan protokol

- ACK kullanarak rdt2.1 ile aynı fonksiyonu görür
- NAK yerine, alıcı en son doğru alınan paket için ACK paket gönderir
 - Alıcı paketin seq numarasını bilmelidir
- Alıcıdaki duplicate ACK : paketin *retransmit edilmesini sağlar*

33/101

rdt2.2: gönderici ve alıcı kısımları



34/101

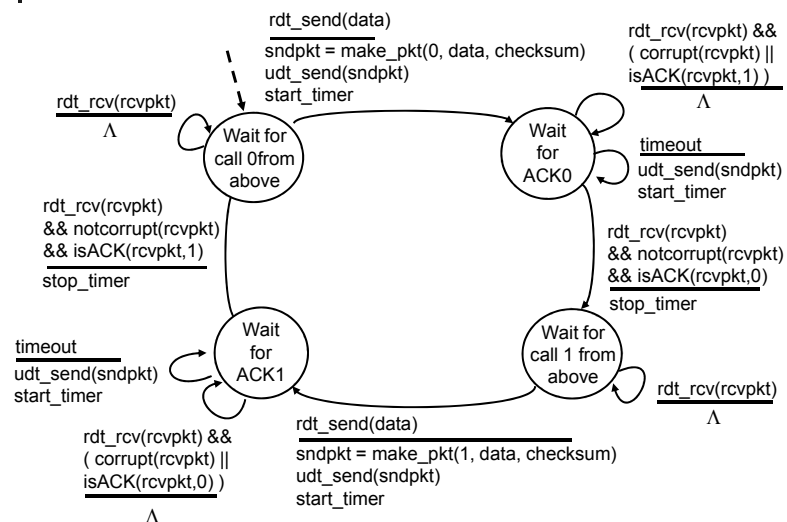
New assumption: underlying channel can also lose packets (data or ACKs)

- checksum, seq. #, ACKs, retransmissions will be of help, but not enough

Approach: sender waits “reasonable” amount of time for ACK



- retransmits if no ACK received in this time
- if pkt (or ACK) just delayed (not lost):
 - retransmission will be duplicate, but use of seq. #'s already handles this
 - receiver must specify seq # of pkt being ACKed
- requires countdown timer

35/101




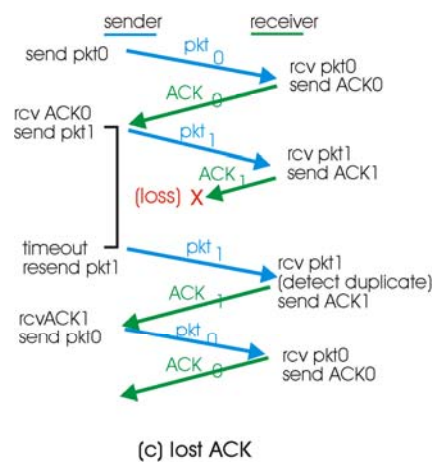
36/101

rdt3.0 in action



37/101

rdt3.0 in action



38/101

Performance of rdt3.0

- rdt3.0 works, but performance stinks
- example: 1 Gbps link, 15 ms e-e prop. delay, 1KB packet:


$$T_{\text{transmit}} = \frac{L \text{ (packet length in bits)}}{R \text{ (transmission rate, bps)}} = \frac{8\text{kb/pkt}}{10^{**9} \text{ b/sec}} = 8 \text{ microsec}$$

$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- U_{sender} : **utilization** – fraction of time sender busy sending
- 1KB pkt every 30 msec -> 33kB/sec thrupt over 1 Gbps link
- network protocol limits use of physical resources!

39/101

rdt3.0: stop-and-wait operation



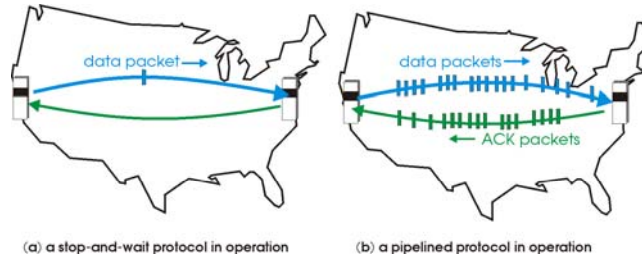
$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

40/101

Pipelined protocols

Pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts


- range of sequence numbers must be increased
- buffering at sender and/or receiver



- Two generic forms of pipelined protocols: *go-Back-N*, *selective repeat*

41/101

Pipelining: increased utilization




42/101

Go-Back-N

Sender:

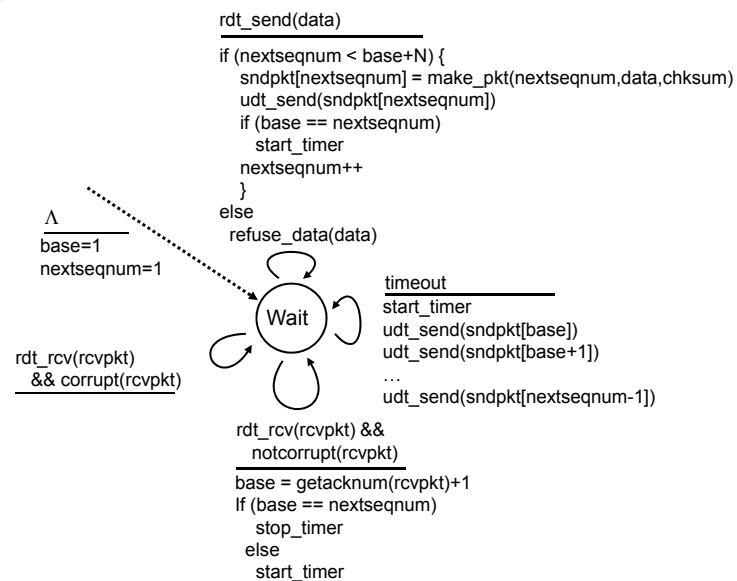
- k-bit seq # in pkt header
- “window” of up to N, consecutive unack’ed pkts allowed



- ACK(n): ACKs all pkts up to, including seq # n - “cumulative ACK”
 - may deceive duplicate ACKs (see receiver)
- timer for each in-flight pkt
- *timeout(n)*: retransmit pkt n and all higher seq # pkts in window


43/101

GRN: sender extended FSM



44/101

GBN: receiver extended FSM




ACK-only: always send ACK for correctly-received pkt with highest *in-order* seq #

- may generate duplicate ACKs
- need only remember **expectedseqnum**
- out-of-order pkt:
 - discard (don't buffer) -> **no receiver buffering!**
 - Re-ACK pkt with highest in-order seq #

45/101

GBN in action



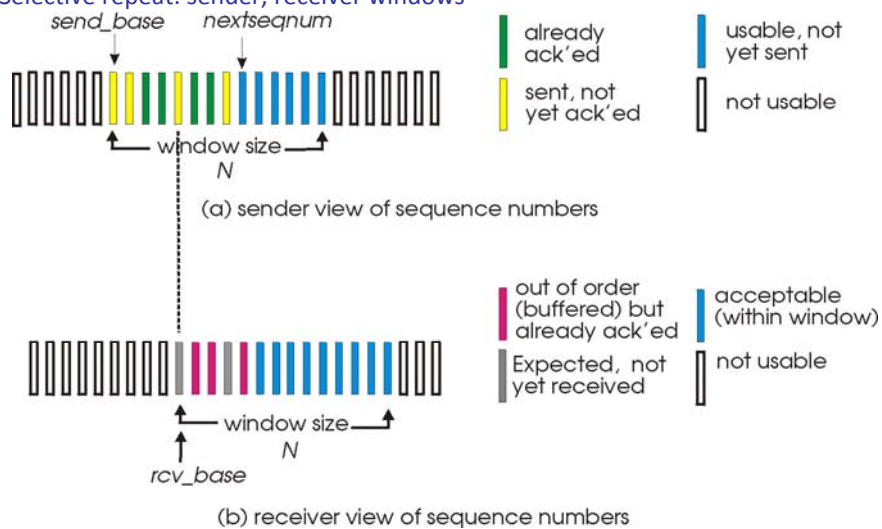
46/101

Selective Repeat

- receiver *individually* acknowledges all correctly received pkts
 - buffers pkts, as needed, for eventual in-order delivery to upper layer
- sender only resends pkts for which ACK not received
 - sender timer for each unACKed pkt
- sender window
 - N consecutive seq #'s
 - again limits seq #'s of sent, unACKed pkts

47/101

Selective repeat: sender, receiver windows



48/101

Selective repeat

sender

data from above :

- if next available seq # in window, send pkt

timeout(n):

- resend pkt n, restart timer

ACK(n) in [sendbase, sendbase+N]:

- mark pkt n as received
- if n smallest unACKed pkt, advance window base to next unACKed seq #

receiver

pkt n in [rcvbase, rcvbase+N-1]

- send ACK(n)
- out-of-order: buffer
- in-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt

pkt n in [rcvbase-N, rcvbase-1]

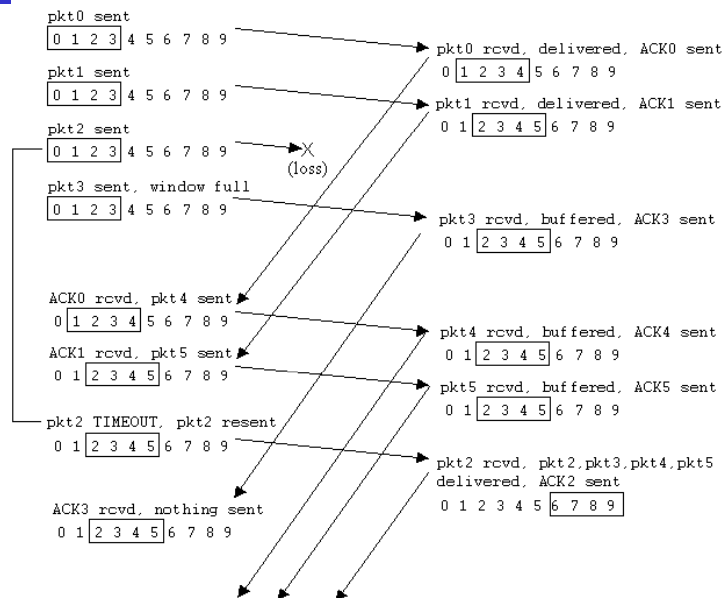
- ACK(n)

otherwise:

- ignore

49/101

Selective repeat in action

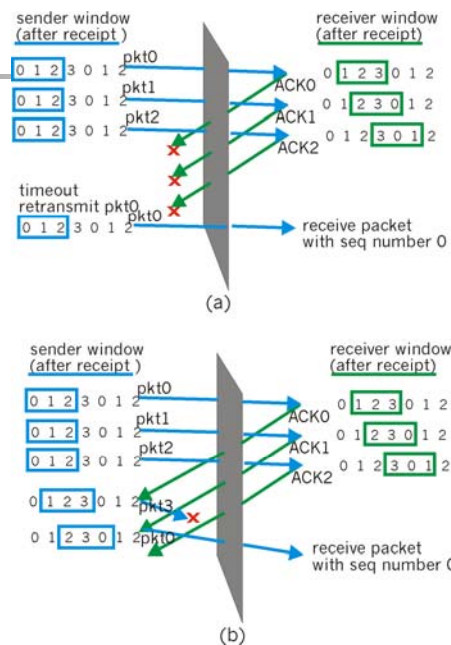


50/101

Selective repeat: dilemma

Example:

- seq #'s: 0, 1, 2, 3
 - window size=3
 - receiver sees no difference in two scenarios!
 - incorrectly passes duplicate data as new in (a)
- Q:** what relationship between seq # size and window size?



51/101

Chapter 3 outline

- 3.1 Transport-layer services
- 3.2 Multiplexing and demultiplexing
- 3.3 Connectionless transport: UDP
- 3.4 Principles of reliable data transfer
- 3.5 Connection-oriented transport: TCP
 - segment structure
 - reliable data transfer
 - flow control
 - connection management
- 3.6 Principles of congestion control
- 3.7 TCP congestion control

52/101

TCP: Overview

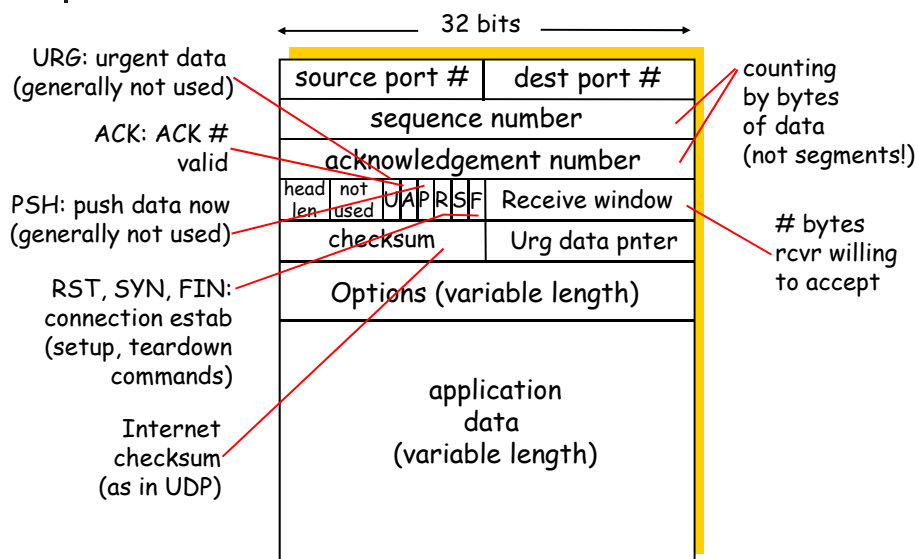
RFCs: 793, 1122, 1323, 2018, 2581

- **point-to-point:**
 - one sender, one receiver
- **reliable, in-order *byte stream*:**
 - no "message boundaries"
- **pipelined:**
 - TCP congestion and flow control set window size
- **send & receive buffers**
- **full duplex data:**
 - bi-directional data flow in same connection
 - MSS: maximum segment size
- **connection-oriented:**
 - handshaking (exchange of control msgs) init's sender, receiver state before data exchange
- **flow controlled:**
 - sender will not overwhelm receiver



53/101

TCP segment structure



54/101

TCP seq. #'s and ACKs

Seq. #'s:

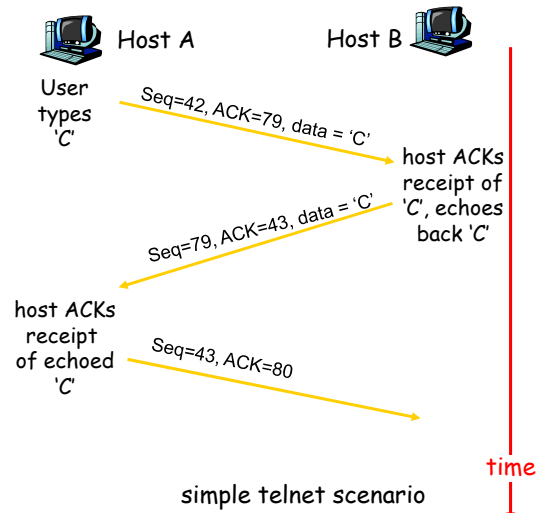
- byte stream
"number" of first byte
in segment's data

ACKs:

- seq # of next byte
expected from other
side
- cumulative ACK

Q: how receiver handles out-
of-order segments

- A: TCP spec doesn't
say, - up to
implementor



55/101

TCP Round Trip Time and Timeout

Q: how to set TCP
timeout value?

- longer than RTT
 - but RTT varies
- too short: premature
timeout
 - unnecessary
retransmissions
- too long: slow reaction to
segment loss

Q: how to estimate RTT?

- SampleRTT:** measured time from
segment transmission until ACK
receipt
 - ignore retransmissions
- SampleRTT** will vary, want
estimated RTT "smoother"
 - average several recent
measurements, not just current
SampleRTT

56/101


TCP Round Trip Time and Timeout

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$

- Exponential weighted moving average
- influence of past sample decreases exponentially fast
- typical value: $\alpha = 0.125$

57/101

Example RTT estimation:



58/101

TCP Round Trip Time and Timeout

Setting the timeout

- **EstimatedRTT** plus “safety margin”
 - large variation in **EstimatedRTT** -> larger safety margin
- first estimate of how much **SampleRTT** deviates from **EstimatedRTT**:

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically, $\beta = 0.25$)

Then set timeout interval:

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$

59/101

Chapter 3 outline

- 3.1 Transport-layer services
- 3.2 Multiplexing and demultiplexing
- 3.3 Connectionless transport: UDP
- 3.4 Principles of reliable data transfer
- 3.5 Connection-oriented transport: TCP
 - segment structure
 - **reliable data transfer**
 - flow control
 - connection management
- 3.6 Principles of congestion control
- 3.7 TCP congestion control

60/101

TCP reliable data transfer

- TCP creates rdt service on top of IP's unreliable service
- Pipelined segments
- Cumulative acks
- TCP uses single retransmission timer
- Retransmissions are triggered by:
 - timeout events
 - duplicate acks
- Initially consider simplified TCP sender:
 - ignore duplicate acks
 - ignore flow control, congestion control

61/101

TCP sender events:

data rcvd from app:

- Create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running (think of timer as for oldest unacked segment)
- expiration interval: TimeoutInterval

timeout:

- retransmit segment that caused timeout
- restart timer

Ack rcvd:

- If acknowledges previously unacked segments
 - update what is known to be acked
 - start timer if there are outstanding segments

62/101

```

NextSeqNum = InitialSeqNum
SendBase = InitialSeqNum

```

```

loop (forever) {
  switch(event)

```

```

    event: data received from application above
           create TCP segment with sequence number NextSeqNum
           if (timer currently not running)
               start timer
           pass segment to IP
           NextSeqNum = NextSeqNum + length(data)

```

```

    event: timer timeout
           retransmit not-yet-acknowledged segment with
             smallest sequence number
           start timer

```

```

    event: ACK received, with ACK field value of y
           if (y > SendBase) {
               SendBase = y
               if (there are currently not-yet-acknowledged segments)
                   start timer
           }

```

```

} /* end of loop forever */

```

TCP sender (simplified)

Comment:

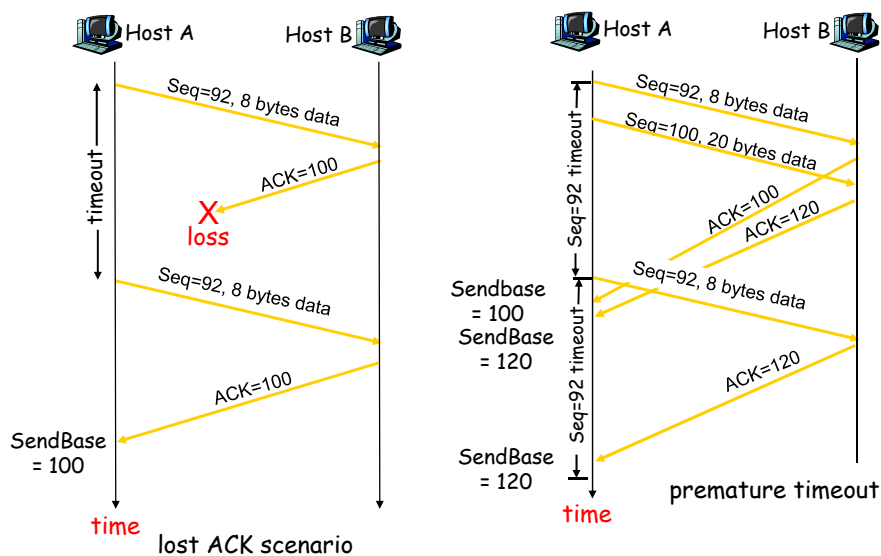
- SendBase-1: last cumulatively ack'd byte

Example:

- SendBase-1 = 71; y = 73, so the rcvr wants 73+ ; y > SendBase, so that new data is acked


63/101

retransmission scenarios



64/101

TCP retransmission scenarios (more)



65/101

TCP ACK generation [RFC 1122, RFC 2581]

Event at Receiver

TCP Receiver action

Arrival of in-order segment with expected seq #. All data up to expected seq # already ACKed

Delayed ACK. Wait up to 500ms for next segment. If no next segment, send ACK

Arrival of in-order segment with expected seq #. One other segment has ACK pending

Immediately send single cumulative ACK, ACKing both in-order segments

Arrival of out-of-order segment higher-than-expected seq. #. Gap detected

Immediately send duplicate ACK, indicating seq. # of next expected byte

Arrival of segment that partially or completely fills gap

Immediate send ACK, provided that segment starts at lower end of gap

66/101

Fast Retransmit

- Time-out period often relatively long:
 - long delay before resending lost packet
- Detect lost segments via duplicate ACKs.
 - Sender often sends many segments back-to-back
 - If segment is lost, there will likely be many duplicate ACKs.
- If sender receives 3 ACKs for the same data, it supposes that segment after ACKed data was lost:
 - **fast retransmit**: resend segment before timer expires

67/101

Fast retransmit algorithm:

```
event: ACK received, with ACK field value of y
  if (y > SendBase) {
    SendBase = y
    if (there are currently not-yet-acknowledged segments)
      start timer
  }
  else {
    increment count of dup ACKs received for y
    if (count of dup ACKs received for y = 3) {
      resend segment with sequence number y
    }
  }
```

a duplicate ACK for
already ACKed segment

fast retransmit

68/101


Chapter 3 outline

- 3.1 Transport-layer services
- 3.2 Multiplexing and demultiplexing
- 3.3 Connectionless transport: UDP
- 3.4 Principles of reliable data transfer
- 3.5 Connection-oriented transport: TCP
 - segment structure
 - reliable data transfer
 - **flow control**
 - connection management
- 3.6 Principles of congestion control
- 3.7 TCP congestion control

69/101

TCP Flow Control

- receive side of TCP connection has a receive buffer:




flow control

sender won't overflow receiver's buffer by transmitting too much, too fast

- speed-matching service: matching the send rate to the receiving app's drain rate
- app process may be slow at reading from buffer

70/101

TCP Flow control: how it works



(Suppose TCP receiver discards out-of-order segments)

- spare room in buffer
- = **RcvWindow**
- = **RcvBuffer - [LastByteRcvd - LastByteRead]**

- Rcvr advertises spare room by including value of **RcvWindow** in segments
- Sender limits unACKed data to **RcvWindow**
 - guarantees receive buffer doesn't overflow

71/101

Chapter 3 outline

- 3.1 Transport-layer services
- 3.2 Multiplexing and demultiplexing
- 3.3 Connectionless transport: UDP
- 3.4 Principles of reliable data transfer
- 3.5 Connection-oriented transport: TCP
 - segment structure
 - reliable data transfer
 - flow control
 - **connection management**
- 3.6 Principles of congestion control
- 3.7 TCP congestion control

72/101

TCP Connection Management

Recall: TCP sender, receiver establish "connection" before exchanging data segments

- initialize TCP variables:
 - seq. #s
 - buffers, flow control info (e.g. **RcvWindow**)
- *client*: connection initiator

```
Socket clientSocket = new
Socket("hostname", "port
number");
```
- *server*: contacted by client

```
Socket connectionSocket =
welcomeSocket.accept();
```

Three way handshake:

Step 1: client host sends TCP SYN segment to server

- specifies initial seq #
- no data

Step 2: server host receives SYN, replies with SYNACK segment

- server allocates buffers
- specifies server initial seq. #

Step 3: client receives SYNACK, replies with ACK segment, which may contain data

73/101


TCP Connection Management (cont.)

Closing a connection:

client closes socket:
`clientSocket.close();`

Step 1: **client** end system sends TCP FIN control segment to server

Step 2: **server** receives FIN, replies with ACK. Closes connection, sends FIN.



74/101


TCP Connection Management (cont.)

Step 3: client receives FIN, replies with ACK.

- Enters "timed wait" - will respond with ACK to received FINs

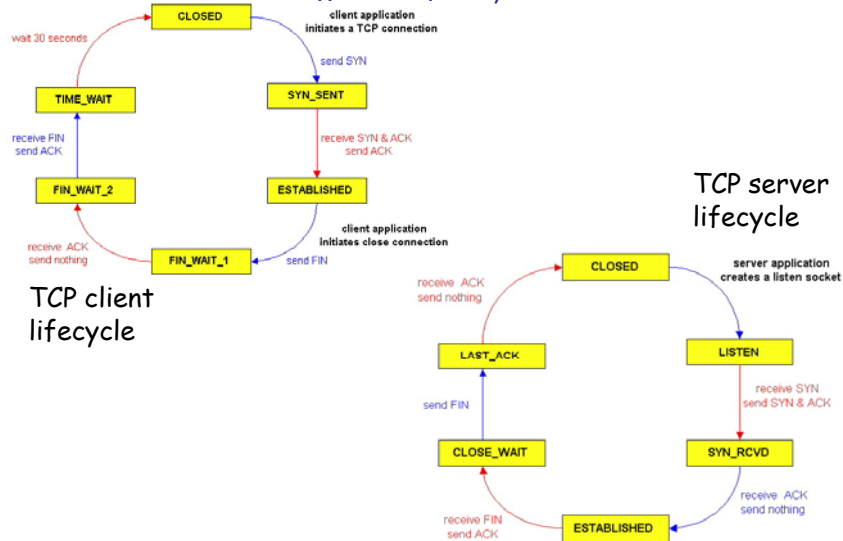
Step 4: server receives ACK. Connection closed.

Note: with small modification, can handle simultaneous FINs.



75/101

TCP Connection Management (cont.)



76/101

Chapter 3 outline

- 3.1 Transport-layer services
- 3.2 Multiplexing and demultiplexing
- 3.3 Connectionless transport: UDP
- 3.4 Principles of reliable data transfer
- 3.5 Connection-oriented transport: TCP
 - segment structure
 - reliable data transfer
 - flow control
 - connection management
- 3.6 Principles of congestion control
- 3.7 TCP congestion control

77/101

Principles of Congestion Control




Congestion:

- informally: “too many sources sending too much data too fast for *network* to handle”
- different from flow control!
- manifestations:
 - lost packets (buffer overflow at routers)
 - long delays (queueing in router buffers)
- a top-10 problem!

78/101

Causes/costs of congestion: scenario 1

- two senders, two receivers
- one router, infinite buffers
- no retransmission




- large delays when congested
- maximum achievable throughput

79/101

Causes/costs of congestion: scenario 2


- one router, *finite* buffers
- sender retransmission of lost packet



80/101

Causes/costs of congestion: scenario 2

- always: $\lambda_{in} = \lambda_{out}$ (goodput)
- “perfect” retransmission only when loss: $\lambda'_{in} > \lambda_{out}$
- retransmission of delayed (not lost) packet makes λ'_{in} larger (than perfect case) for same λ_{out}



“costs” of congestion:


- more work (retrans) for given “goodput”
- unneeded retransmissions: link carries multiple copies of pkt

81/101

Causes/costs of congestion: scenario 3


- four senders
- multihop paths
- timeout/retransmit

Q: what happens as λ_{in} and λ'_{in} increase?



82/101

Causes/costs of congestion: scenario 3



Another “cost” of congestion:

- when packet dropped, any “upstream transmission capacity used for that packet was wasted!

83/101

Approaches towards congestion control

Two broad approaches towards congestion control:

End-end congestion control:

- no explicit feedback from network
- congestion inferred from end-system observed loss, delay
- approach taken by TCP

Network-assisted congestion control:

- routers provide feedback to end systems
 - single bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM)
 - explicit rate sender should send at

84/101

Case study: ATM ABR congestion control

ABR: available bit rate:


- “elastic service”
- if sender’s path “underloaded”:
 - sender should use available bandwidth
- if sender’s path congested:
 - sender throttled to minimum guaranteed rate

RM (resource management) cells:

- sent by sender, interspersed with data cells
- bits in RM cell set by switches (“network-assisted”)
 - **NI bit**: no increase in rate (mild congestion)
 - **CI bit**: congestion indication
- RM cells returned to sender by receiver, with bits intact

85/101

Case study: ATM ABR congestion control



- two-byte ER (explicit rate) field in RM cell
 - congested switch may lower ER value in cell
 - sender’ send rate thus minimum supportable rate on path
- EFCI bit in data cells: set to 1 in congested switch
 - if data cell preceding RM cell has EFCI set, sender sets CI bit in returned RM cell

86/101

Chapter 3 outline

- 3.1 Transport-layer services
- 3.2 Multiplexing and demultiplexing
- 3.3 Connectionless transport: UDP
- 3.4 Principles of reliable data transfer
- 3.5 Connection-oriented transport: TCP
 - segment structure
 - reliable data transfer
 - flow control
 - connection management
- 3.6 Principles of congestion control
- 3.7 TCP congestion control

87/101

TCP Congestion Control

- end-end control (no network assistance)
 - sender limits transmission:
 $\text{LastByteSent} - \text{LastByteAcked} \leq \text{CongWin}$
 - Roughly,


$$\text{rate} = \frac{\text{CongWin}}{\text{RTT}} \text{ Bytes/sec}$$
 - **CongWin** is dynamic, function of perceived network congestion
- How does sender perceive congestion?
- loss event = timeout or 3 duplicate acks
 - TCP sender reduces rate (**CongWin**) after loss event
- three mechanisms:
- AIMD
 - slow start
 - conservative after timeout events

88/101

TCP AIMD

multiplicative decrease: cut **CongWin** in half after loss event

additive increase: increase **CongWin** by 1 MSS every RTT in the absence of loss events: *probing*



Long-lived TCP connection

89/101


TCP Slow Start

- When connection begins, **CongWin** = 1 MSS
 - Example: MSS = 500 bytes & RTT = 200 msec
 - initial rate = 20 kbps
- available bandwidth may be $\gg \text{MSS}/\text{RTT}$
 - desirable to quickly ramp up to respectable rate
- When connection begins, increase rate exponentially fast until first loss event

90/101

TCP Slow Start (more)

- When connection begins, increase rate exponentially until first loss event:
 - double **CongWin** every RTT
 - done by incrementing **CongWin** for every ACK received
- **Summary:** initial rate is slow but ramps up exponentially fast



91/101

Refinement

- After 3 dup ACKs:
 - **CongWin** is cut in half
 - window then grows linearly
- **But** after timeout event:
 - **CongWin** instead set to 1 MSS;
 - window then grows exponentially
 - to a threshold, then grows linearly

Philosophy:


- 3 dup ACKs indicates network capable of delivering some segments
- timeout before 3 dup ACKs is "more alarming"

92/101

Refinement (more)

Q: When should the exponential increase switch to linear?

A: When **CongWin** get to 1/2 of its value before timeout.



Implementation:

- Variable Threshold
- At loss event, Threshold is set to 1/2 of CongWin just before loss event

93/101

Summary: TCP Congestion Control

- When **CongWin** is below **Threshold**, sender in **slow-start** phase, window grows exponentially.
- When **CongWin** is above **Threshold**, sender is in **congestion-avoidance** phase, window grows linearly.
- When a **triple duplicate ACK** occurs, **Threshold** set to **CongWin/2** and **CongWin** set to **Threshold**.
- When **timeout** occurs, **Threshold** set to **CongWin/2** and **CongWin** is set to 1 MSS.

94/101

TCP sender congestion control

Event	State	TCP Sender Action	Commentary
ACK receipt for previously unacked data	Slow Start (SS)	$\text{CongWin} = \text{CongWin} + \text{MSS}$, If ($\text{CongWin} > \text{Threshold}$) set state to "Congestion Avoidance"	Resulting in a doubling of CongWin every RTT
ACK receipt for previously unacked data	Congestion Avoidance (CA)	$\text{CongWin} = \text{CongWin} + \text{MSS} * (\text{MSS} / \text{CongWin})$	Additive increase, resulting in increase of CongWin by 1 MSS every RTT
Loss event detected by triple duplicate ACK	SS or CA	$\text{Threshold} = \text{CongWin} / 2$, $\text{CongWin} = \text{Threshold}$, Set state to "Congestion Avoidance"	Fast recovery, implementing multiplicative decrease. CongWin will not drop below 1 MSS.
Timeout	SS or CA	$\text{Threshold} = \text{CongWin} / 2$, $\text{CongWin} = 1 \text{ MSS}$, Set state to "Slow Start"	Enter slow start
Duplicate ACK	SS or CA	Increment duplicate ACK count for segment being acked	CongWin and Threshold not changed

95/101

TCP throughput

- What's the average throughput of TCP as a function of window size and RTT?
 - Ignore slow start
- Let W be the window size when loss occurs.
- When window is W , throughput is W/RTT
- Just after loss, window drops to $W/2$, throughput to $W/2\text{RTT}$.
- Average throughput: $.75 W/\text{RTT}$

96/101

TCP Futures


- Example: 1500 byte segments, 100ms RTT, want 10 Gbps throughput
- Requires window size $W = 83,333$ in-flight segments
- Throughput in terms of loss rate:

- $$\frac{1.22 \cdot MSS}{RTT \sqrt{L}}$$
- $\rightarrow L = 2 \cdot 10^{-10}$ *Wow*
 - New versions of TCP for high-speed needed!

97/101

TCP Fairness

Fairness goal: if K TCP sessions share same bottleneck link of bandwidth R , each should have average rate of R/K




98/101

Why is TCP fair?

Two competing sessions:

- Additive increase gives slope of 1, as throughput increases
- multiplicative decrease decreases throughput proportionally



99/101

Fairness (more)

Fairness and UDP

- Multimedia apps often do not use TCP
 - do not want rate throttled by congestion control
- Instead use UDP:
 - pump audio/video at constant rate, tolerate packet loss
- Research area: TCP friendly

Fairness and parallel TCP connections

- nothing prevents app from opening parallel cncctions between 2 hosts.
- Web browsers do this
- Example: link of rate R supporting 9 cncctions;
 - new app asks for 1 TCP, gets rate $R/10$
 - new app asks for 11 TCPs, gets $R/2$!

100/101

Delay modeling

Q: How long does it take to receive an object from a Web server after sending a request?

Ignoring congestion, delay is influenced by:

- TCP connection establishment
- data transmission delay
- slow start

Notation, assumptions:

- Assume one link between client and server of rate R
- S : MSS (bits)
- O : object size (bits)
- no retransmissions (no loss, no corruption)

Window size:

- First assume: fixed congestion window, W segments
- Then dynamic window, modeling slow start


101/101

Fixed congestion window (1)

First case:

$WS/R > RTT + S/R$: ACK for first segment in window returns before window's worth of data sent

$$\text{delay} = 2RTT + O/R$$




102/101

Fixed congestion window (2)

Second case:

- $WS/R < RTT + S/R$: wait for ACK after sending window's worth of data sent

$$\text{delay} = 2RTT + O/R + (K-1)[S/R + RTT - WS/R]$$



103/101

TCP Delay Modeling: Slow Start (1)

Now suppose window grows according to slow start

Will show that the delay for one object is:

$$\text{Latency} = 2RTT + \frac{O}{R} + P \left[RTT + \frac{S}{R} \right] - (2^P - 1) \frac{S}{R}$$

where P is the number of times TCP idles at server:

$$P = \min\{Q, K-1\}$$

- where Q is the number of times the server idles if the object were of infinite size.
- and K is the number of windows that cover the object.

104/101

TCP Delay Modeling: Slow Start (2)

Delay components:

- 2 RTT for connection estab and request
- O/R to transmit object
- time server idles due to slow start


Server idles:

$$P = \min\{K-1, Q\} \text{ times}$$

Example:

- $O/S = 15$ segments
- $K = 4$ windows
- $Q = 2$
- $P = \min\{K-1, Q\} = 2$

Server idles $P=2$ times



105/101


TCP Delay Modeling (3)

$\frac{S}{R} + RTT$ = time from when server starts to send segment
until server receives acknowledgement

$$2^{k-1} \frac{S}{R} = \text{time to transmit the } k\text{th window}$$

$$\left[\frac{S}{R} + RTT - 2^{k-1} \frac{S}{R} \right]^+ = \text{idle time after the } k\text{th window}$$

$$\begin{aligned} \text{delay} &= \frac{O}{R} + 2RTT + \sum_{p=1}^P \text{idleTime}_p \\ &= \frac{O}{R} + 2RTT + \sum_{k=1}^P \left[\frac{S}{R} + RTT - 2^{k-1} \frac{S}{R} \right] \\ &= \frac{O}{R} + 2RTT + P \left[RTT + \frac{S}{R} \right] - (2^P - 1) \frac{S}{R} \end{aligned}$$



106/101

TCP Delay Modeling (4)

Recall K = number of windows that cover object

How do we calculate K ?

$$\begin{aligned} K &= \min\{k : 2^0 S + 2^1 S + \dots + 2^{k-1} S \geq O\} \\ &= \min\{k : 2^0 + 2^1 + \dots + 2^{k-1} \geq O/S\} \\ &= \min\{k : 2^k - 1 \geq \frac{O}{S}\} \\ &= \min\{k : k \geq \log_2(\frac{O}{S} + 1)\} \\ &= \left\lceil \log_2(\frac{O}{S} + 1) \right\rceil \end{aligned}$$

Calculation of Q , number of idles for infinite-size object, is similar (see HW).

107/101


HTTP Modeling

- Assume Web page consists of:
 - 1 base HTML page (of size O bits)
 - M images (each of size O bits)
- Non-persistent HTTP:
 - $M+1$ TCP connections in series
 - Response time = $(M+1)O/R + (M+1)2RTT + \text{sum of idle times}$
- Persistent HTTP:
 - 2 RTT to request and receive base HTML file
 - 1 RTT to request and receive M images
 - Response time = $(M+1)O/R + 3RTT + \text{sum of idle times}$
- Non-persistent HTTP with X parallel connections
 - Suppose M/X integer.
 - 1 TCP connection for base file
 - M/X sets of parallel connections for images.
 - Response time = $(M+1)O/R + (M/X + 1)2RTT + \text{sum of idle times}$

108/101

HTTP Response time (in seconds)

RTT = 100 msec, O = 5 Kbytes, M=10 and X=5




For low bandwidth, connection & response time dominated by transmission time.

Persistent connections only give minor improvement over parallel connections.

109/101

HTTP Response time (in seconds)

RTT = 1 sec, O = 5 Kbytes, M=10 and X=5



For larger RTT, response time dominated by TCP establishment & slow start delays. Persistent connections now give important improvement: particularly in high delay-bandwidth networks.

110/101

Chapter 3: Summary

- principles behind transport layer services:
 - multiplexing, demultiplexing
 - reliable data transfer
 - flow control
 - congestion control
 - instantiation and implementation in the Internet
 - UDP
 - TCP
- Next:
- leaving the network “edge” (application, transport layers)
 - into the network “core”

111/101